Setting up a web server load balancer on a computer cluster

A draft tutorial by Konstantinos Delistavrou

April 2018

We need at least 3 nodes to have a cluster.

We will use 2 “real” Raspberry Pis and a “virtual” Raspberry Pi as a virtual machine on the laptop.

Part 1: Setup virtual machine

Now, let’s setup the virtual machine.

Go to https://www.virtualbox.org/wiki/Downloads and download VirtualBox for Windows hosts.

Install it.

Go to https://www.raspberrypi.org/downloads/raspberry-pi-desktop/ and download “The Raspberry Pi Desktop OS for PC and Mac - based on Debian Stretch”.

On VirtualBox create a virtual machine (button New).

Then create a virtual machine like in the pictures.

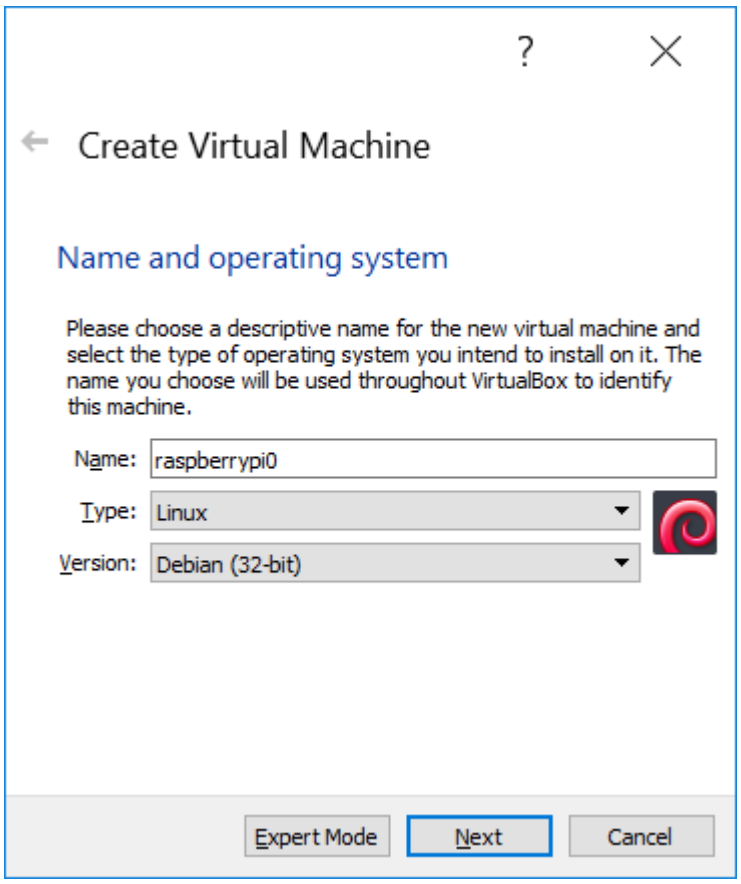

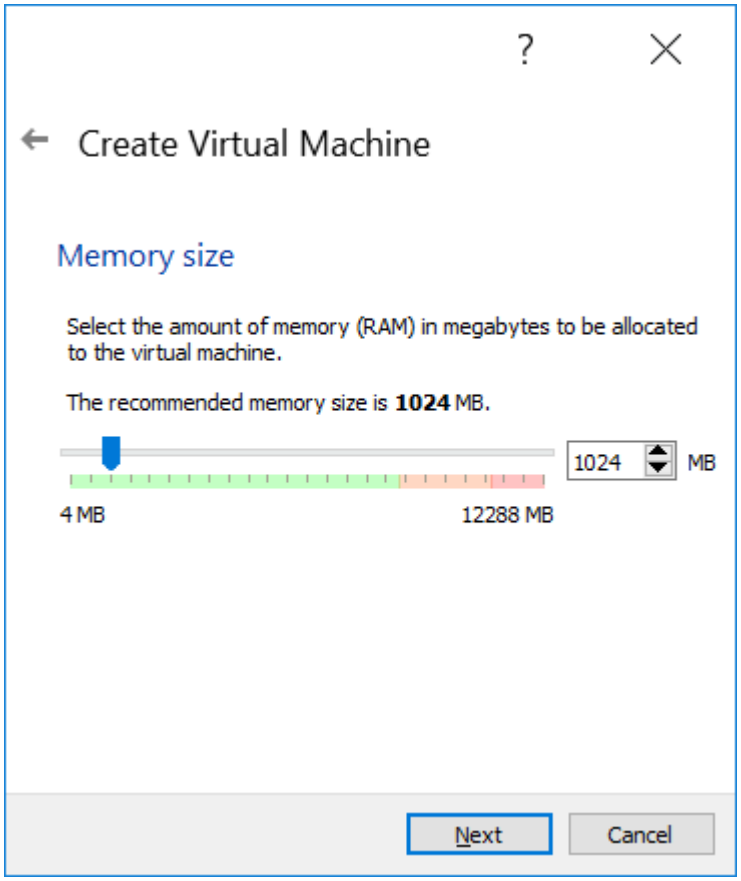
Create Virtual Machine
Memory size
Select the amount of memory (RAM) in megabytes to be allocated to the virtual machine.
The recommended memory size is 1024 MB.
1024
MB
4 MB
12288 MB
Next
Cancel

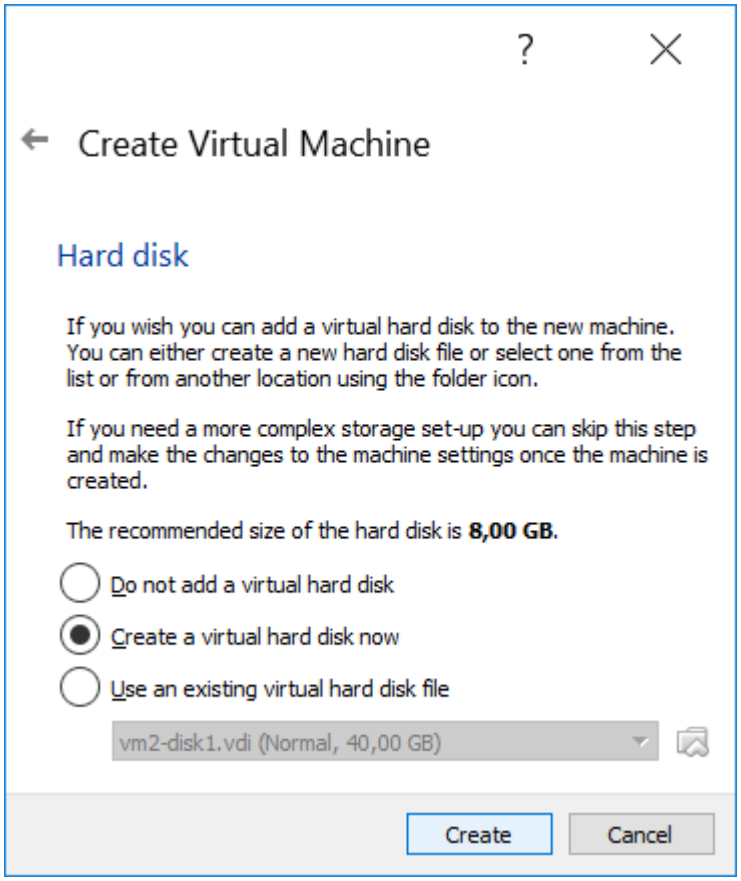
Create Virtual Machine
Hard disk
If you wish you can add a virtual hard disk to the new machine. You can either create a new hard disk file or select one from the list or from another location using the folder icon.
If you need a more complex storage set-up you can skip this step and make the changes to the machine settings once the machine is created.
The recommended size of the hard disk is 8,00 GB.
Do not add a virtual hard disk
Create a virtual hard disk now
Use an existing virtual hard disk file
vm2-disk1.vdi (Normal, 40,00 GB)
Create
Cancel

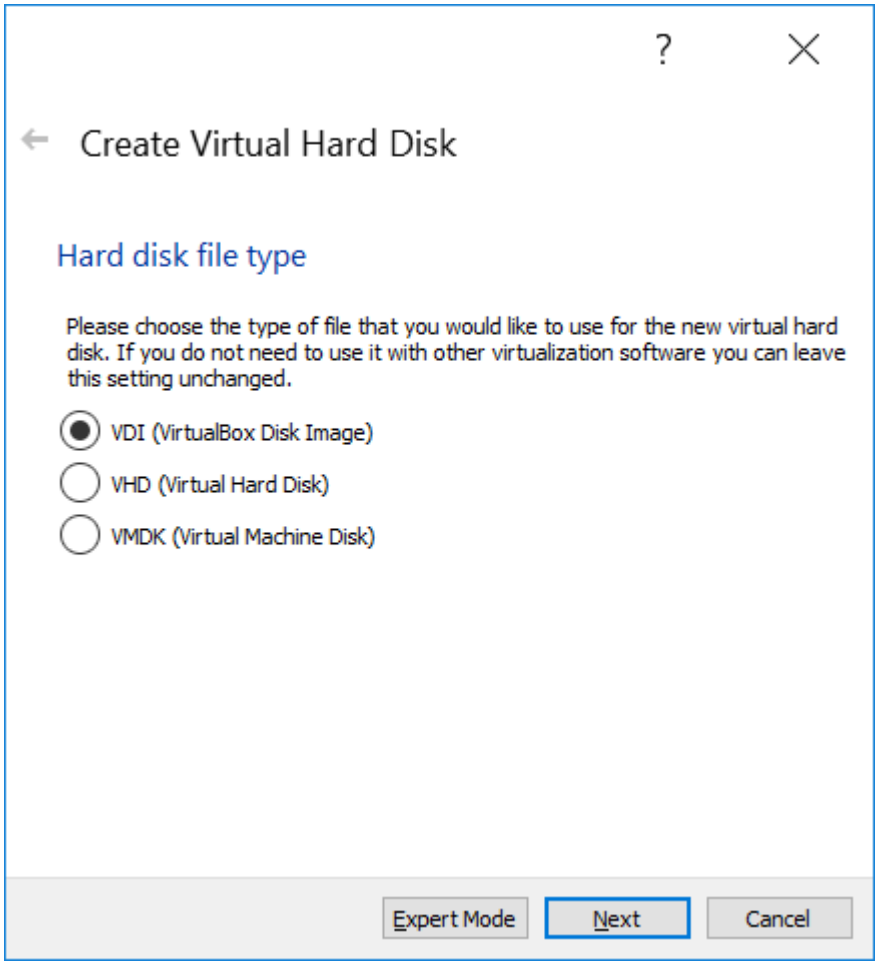
Create Virtual Hard Disk
Hard disk file type
Please choose the type of file that you would like to use for the new virtual hard disk. If you do not need to use it with other virtualization software you can leave this setting unchanged.
VDI (VirtualBox Disk Image)
VHD (Virtual Hard Disk)
VMDK (Virtual Machine Disk)
Expert Mode
Next
Cancel

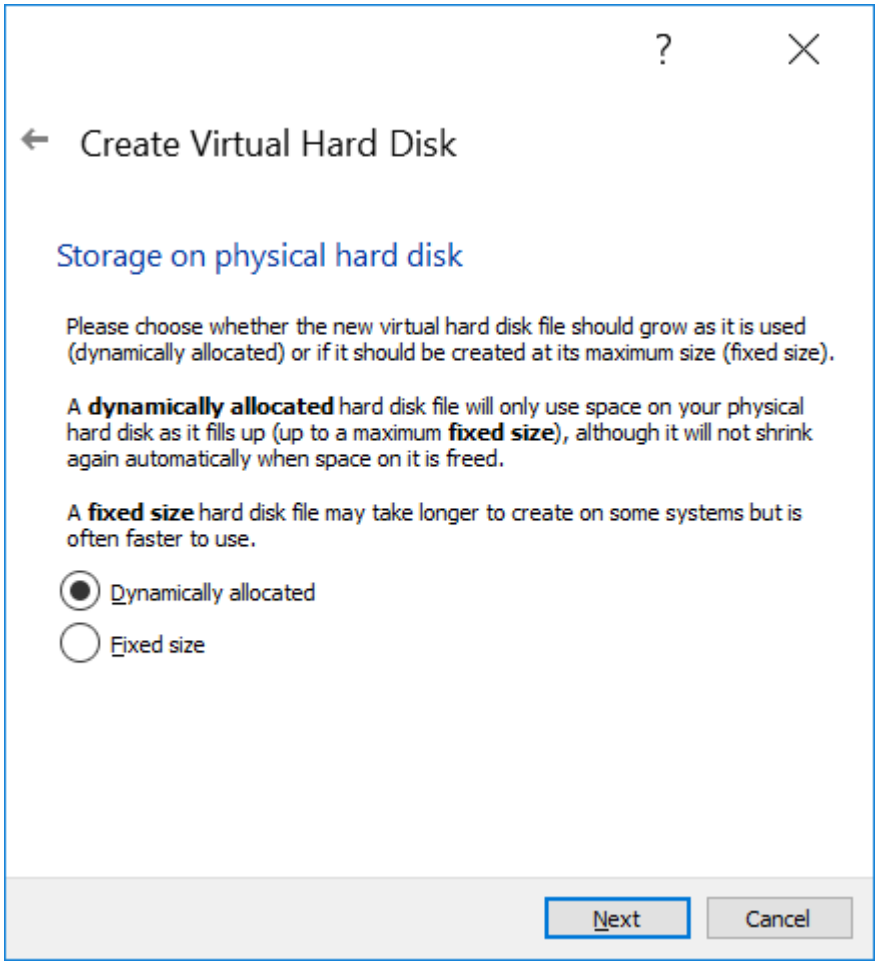
Create Virtual Hard Disk
Storage on physical hard disk
Please choose whether the new virtual hard disk file should grow as it is used (dynamically allocated) or if it should be created at its maximum size (fixed size).
A dynamically allocated hard disk file will only use space on your physical hard disk as it fills up (up to a maximum fixed size), although it will not shrink again automatically when space on it is freed.
A fixed size hard disk file may take longer to create on some systems but is often faster to use.
Dynamically allocated
Fixed size
Next
Cancel

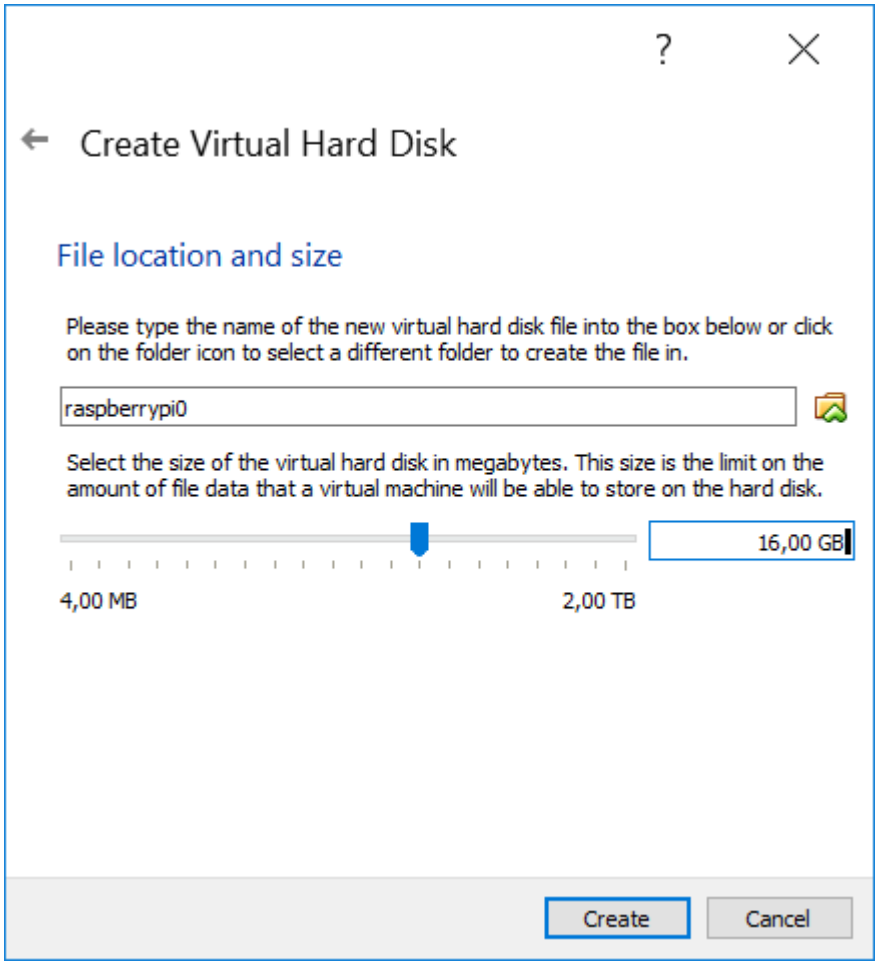

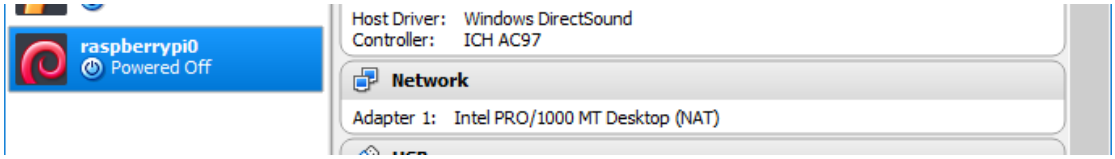

Right click the raspberrypi0 virtual machine and select Settings.

Go to Network tab.

Select Attached to Bridged Adapter.

On Name select the laptop's Wireless network interface (This is important!).

On Adapter choose PCnet-Fast III.

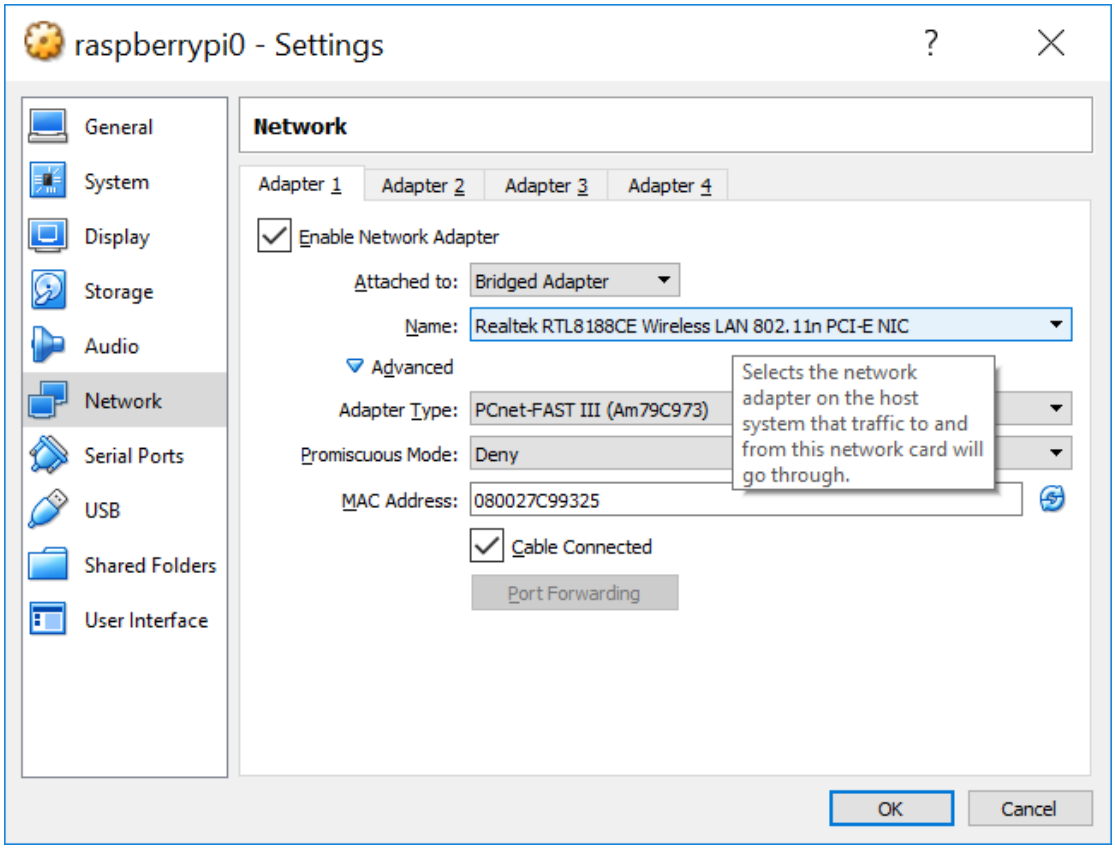

Go to Storage tab.

Select Empty disk and on the right panel click the disk icon.

Select Choose Virtual Optical Disk File.

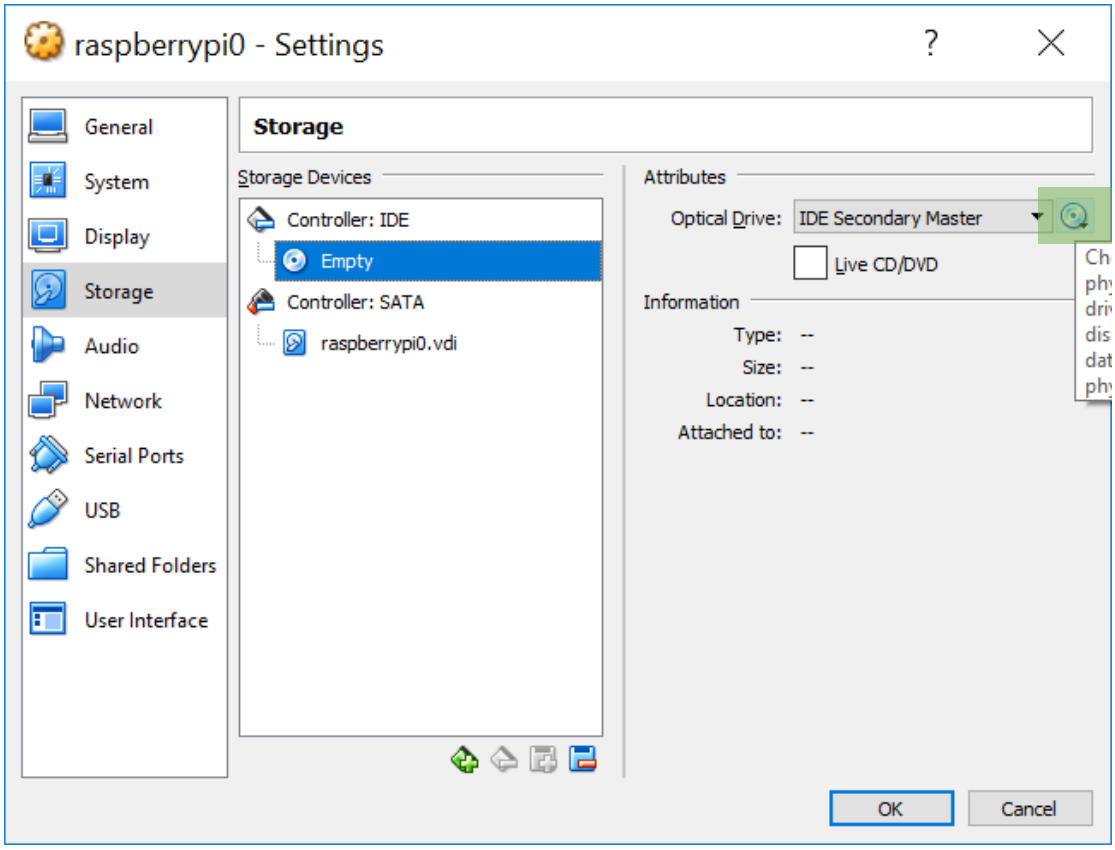

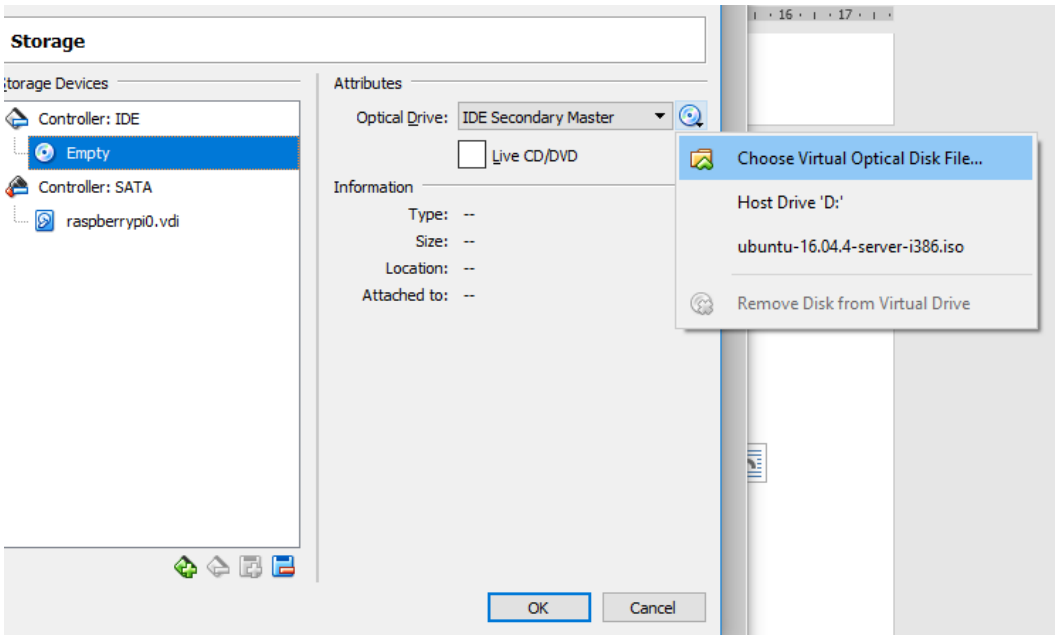

Select the image “2017-11-16-rpd-x86-stretch.iso” downloaded earlier.

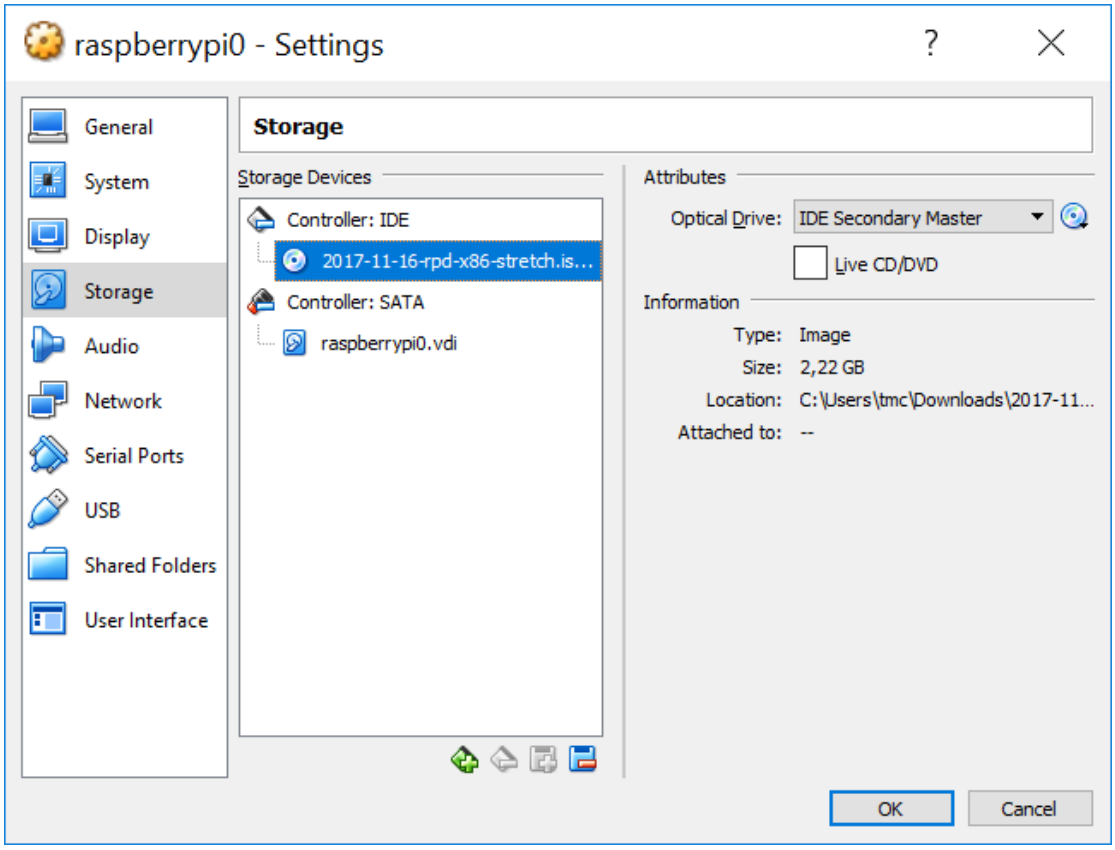

Press OK.

Start the virtual machine.

Choose Graphical Install.

Complete the installation using the default options (When asked to write changes to disk, and to write GRUB, answer Yes!).

See an example in this tutorial https://www.techrepublic.com/article/how-to-run-raspberry-pis-raspbian-os-on-a-pc/

At the end you will have a Raspberry Pi as a virtual machine.

Part 2: Networking

Connect 1st Raspberry Pi on the HDMI monitor, keyboard, mouse.

Power it up.

Open a command terminal.

Give **`ifconfig`** and write to a paper its IP address (mine is 192.168.0.103).

```
wlan0: flags=4163<UP,BROADCAST,RUNNING,MULTICAST>  mtu 1500
        inet 192.168.0.103  netmask 255.255.255.0  broadcast 192.168.0.255
        inet6 fe80::d502:a6a0:df90:63cc  prefixlen 64  scopeid 0x20<link>
        ether b8:27:eb:ba:71:33  txqueuelen 1000  (Ethernet)
        RX packets 302  bytes 27951 (27.2 KiB)
        RX errors 0  dropped 0  overruns 0  frame 0
        TX packets 151  bytes 26561 (25.9 KiB)
        TX errors 0  dropped 0 overruns 0  carrier 0  collisions 0
```

The default username is pi.

The default password is raspberry.

On the terminal type **`sudo nano /etc/hostname`**.

Change the name to **`raspberrypi1`**.

Press Ctrl + O to save.

Press Ctrl + X to exit nano.

```
  GNU nano 2.7.4                File: /etc/hostname

raspberrypi1


                              [ Wrote 1 line ]
^G Get Help  ^O Write Out ^W Where Is  ^K Cut Text  ^J Justify   ^C Cur Pos
^X Exit      ^R Read File ^\ Replace   ^U Uncut Text^T To Spell  ^_ Go To Line
```

Type **`sudo reboot`** to reboot.

When it boots open a terminal again and type **`ifconfig`** to get its IP. Usually is the same.

Connect 2nd Raspberry Pi on the HDMI monitor, keyboard, mouse.

Do the same steps but name it raspberrypi2.

Start the raspberrypi0 virtual machine.

Do the same steps but name it raspberrypi0.

After reboot usually the Raspberry Pis don't change IP address. If you are not lucky and they change, use Angry IP scanner to find their addresses.

Go to http://angryip.org/download/#windows, download and install it.

Scan the Wi-Fi LAN range (e.g. from 192.168.0.1 to 192.168.0.254) for all IP addresses (.1 or .254 is usually the router). Click on the Ping column and Sort by Ping.

You should see a list like the following. Mark Raspberry Pis addresses.

Do this every time you lose the IPs of the Raspberry Pis.

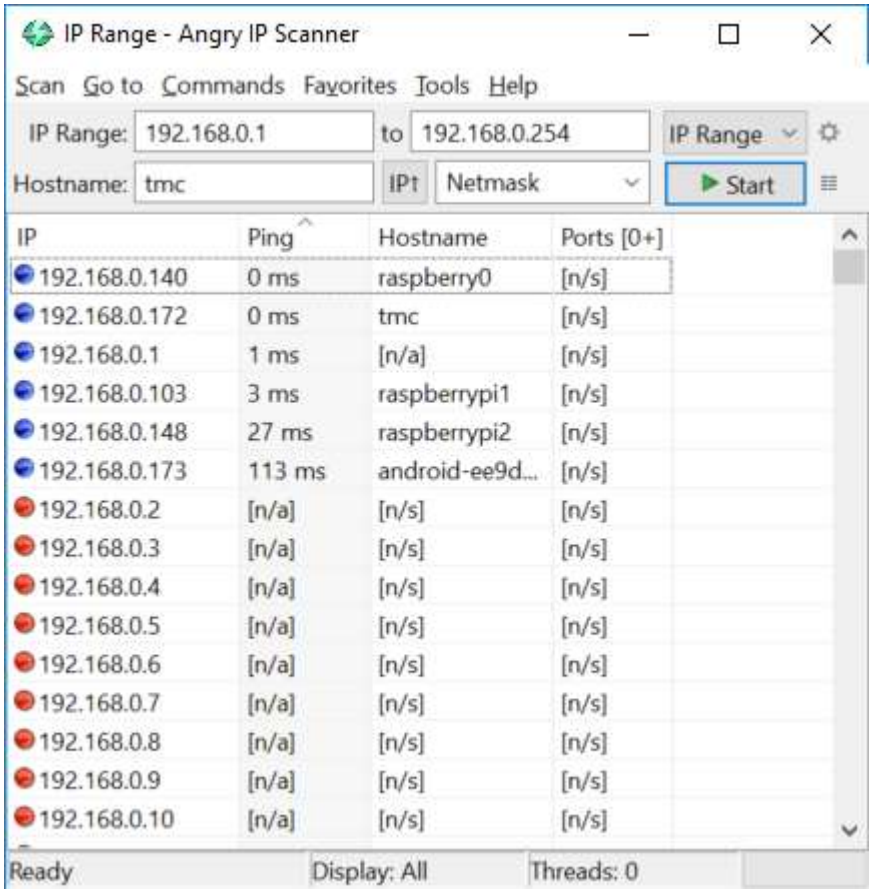

After these first steps you will have 2 Raspberry Pis connected to the Wi-Fi router, each one with a different IP.

Sometimes, after restart, the network interface doesn't work.

In this case it is better to set a static IP address to the Raspberry Pi.

Edit the appropriate file by typing **`sudo nano /etc/dhcpcd.conf`**

Find the line "# Example static IP configuration:" and change the following lines as in the example of the next picture. Also remove the # character, so they are not a comment.

```
# Example static IP configuration:
interface eth0
static ip_address=192.168.0.10/24
#static ip6_address=fd51:42f8:caae:d92e::ff/64
static routers=192.168.0.1
static domain_name_servers=192.168.0.1 8.8.8.8
```

Where …

- "static ip_address=" type a unique IP address for the Wi-Fi LAN subnet,
  prefer numbers after .200, like 192.168.0.201, etc.
- "static routers" is the address of the router.
- "static domain_name_servers" is the address of the router
  and 8.8.8.8 (Google's public DNS server)

Leave the line #static ip6_address commented (don't delete # in front of it).

Reboot the Raspberry Pi to apply changes by typing **`sudo reboot`**

Do this for each Raspberry Pi having problem, but each time with a different IP address (the router and domain name servers must be the same).

Part 3: Remote terminal

It is more convenient to work with a remote terminal, except if you have two HDMI monitors and mouse keyboard sets.

Otherwise you may enable SSH and connect to Raspberry Pis from the laptop.

In this case open a terminal on each Raspberry Pi you want to access remotely.

Type **`sudo raspi-config`**

Use the keyboard's arrow keys, TAB and Enter to browse the menu.

Select 5 Interfacing Options.

Select P2 SSH.

Select <Yes> to the question about enabling the SSH server.

Select Finish.

Reboot Raspberry Pi

On the laptop, download and install PuTTY from this URL
https://www.chiark.greenend.org.uk/~sgtatham/putty/latest.html

Install it.

Run PuTTY and just type the Raspberry Pi's IP address. Press Open.

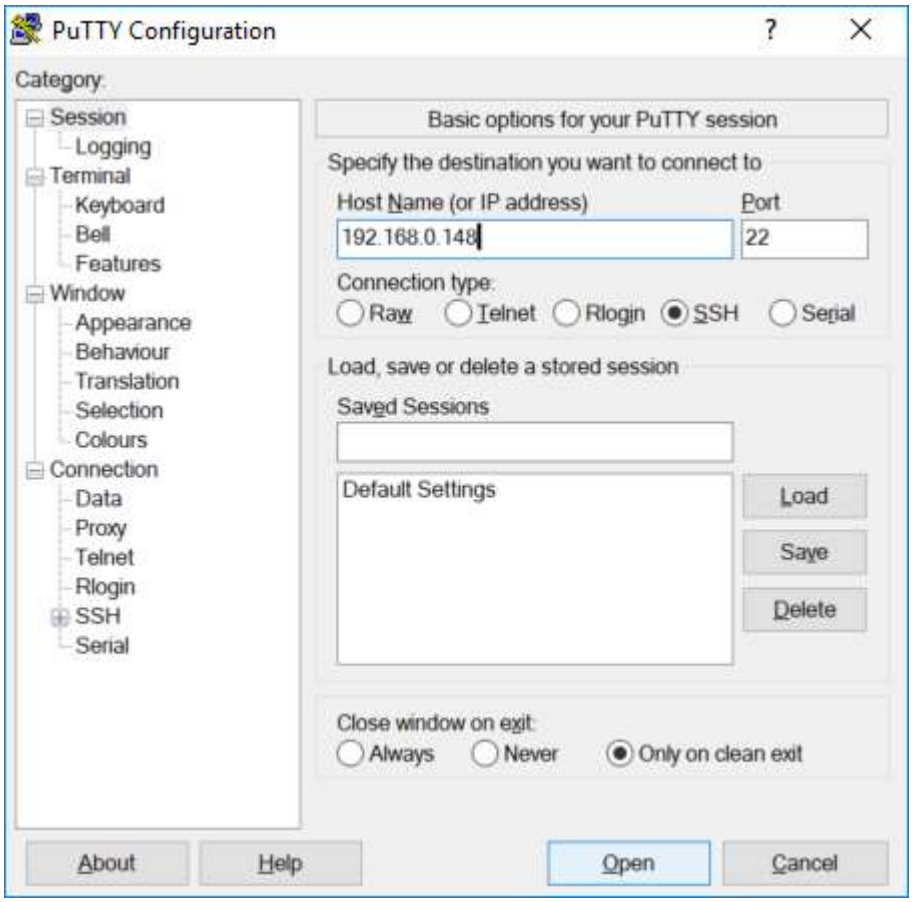

Reply Yes if you are asked about a certificate.

Then login with the username and password (usually pi, raspberry).

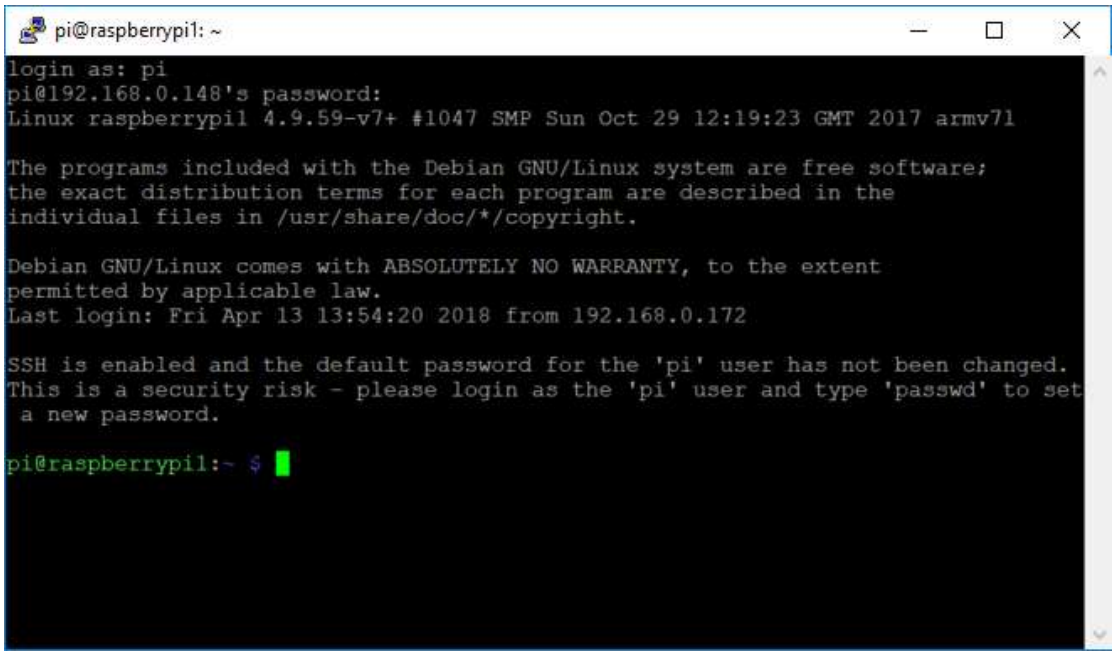

While using PuTTY you can Copy text by selecting it and left clicking on it. We can also Paste text by right clicking on the PuTTY window.

Part 4: Web Server

First update all the Raspberry Pis (0 virtual machine, 1 real, 2 real)

Type (on each raspberry)
`sudo apt-get update`

Then install (on each raspberry) the Apache web server
`sudo apt-get install apache2`

Open a web browser and type the IP address of each Raspberry Pi.

We should see that apache is running, like in the following pictures.

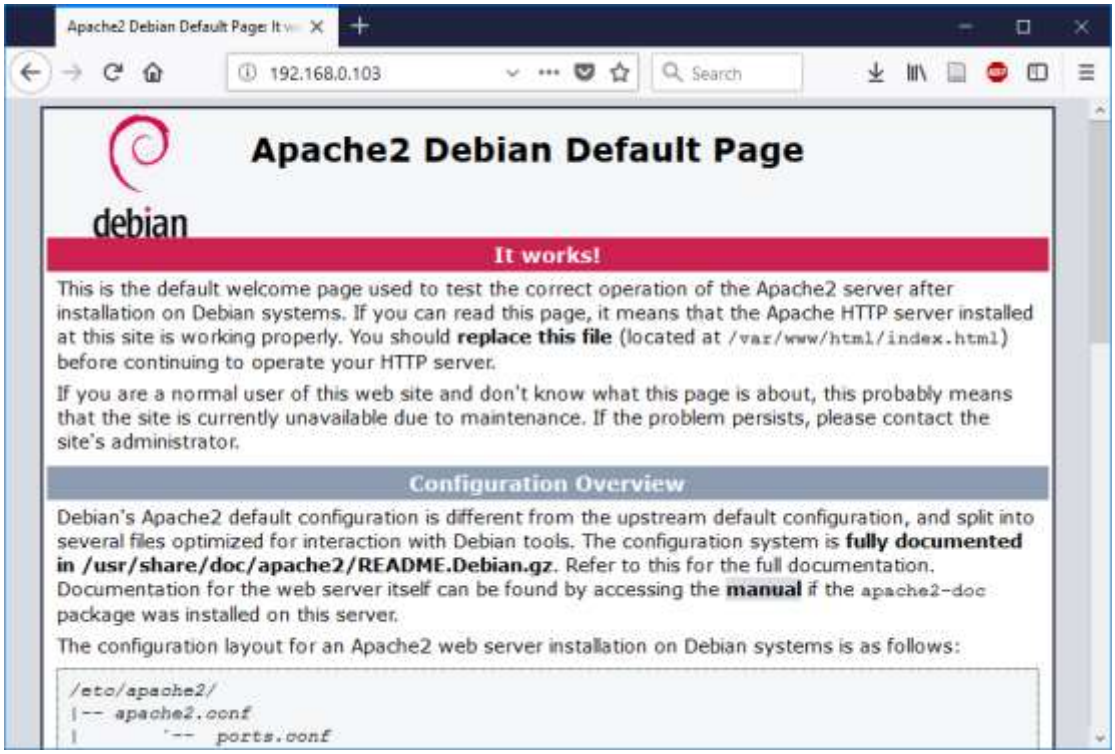

To know which Raspberry Pi web server is responding, you should add one message to each default web page.

Go to each Raspberry Pi and do the following…

Open a terminal and type

`cd /var/www/html`

sudo nano index.html

Find the line <body> …

and bellow it type a new line…

<h2>This is RaspberryPi1</h2>

like in the picture.

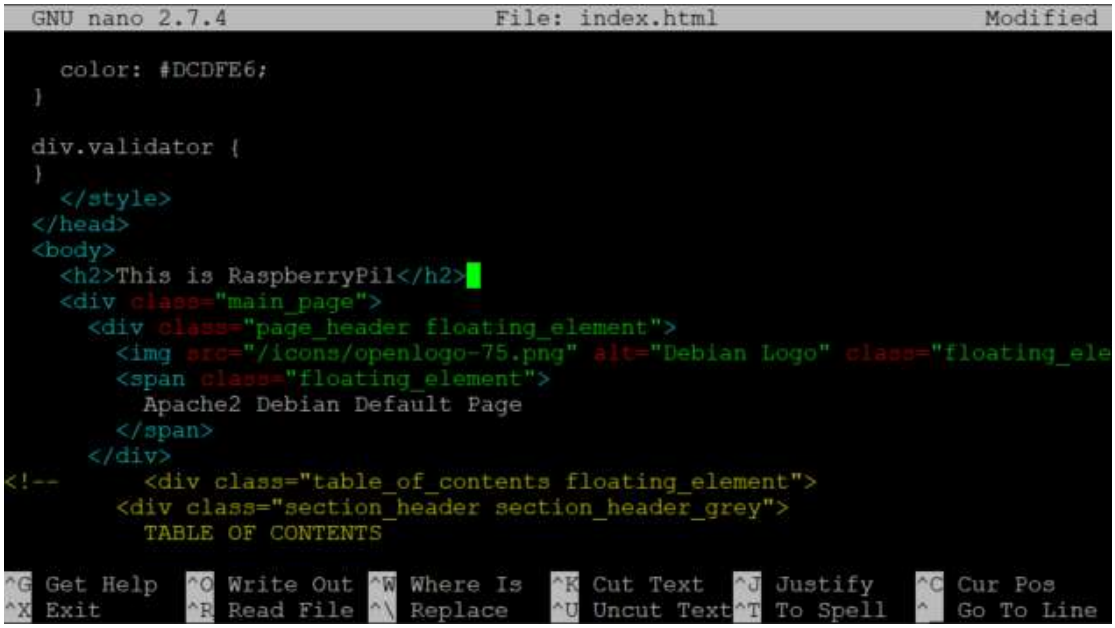

Press Ctrl + O to save and Ctrl + X to quit nano.

Go to the browser and type the Raspberry Pi's IP address. We should something like the following.

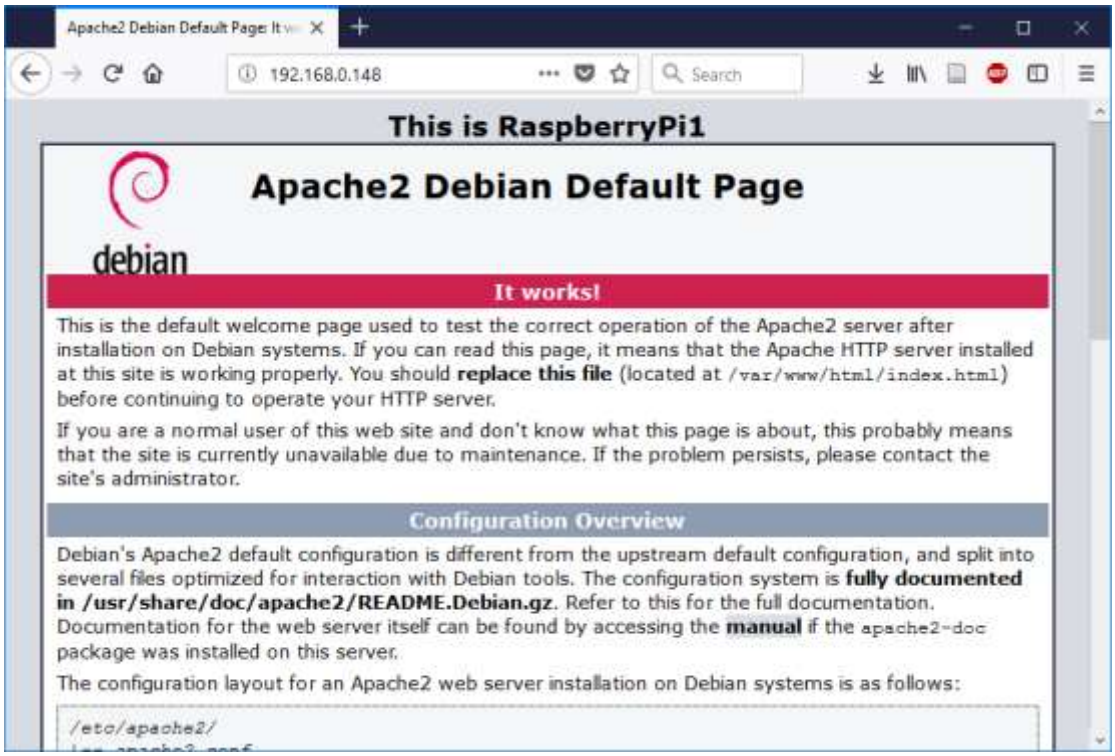

Thus, you can recognize which Raspberry Pi is serving the web page.

Do this for all the Raspberry Pis of the cluster, changing the number according to the number of each Raspberry Pi.

Change the port of each web server.

On each Raspberry Pi edit the ports configuration file and change the port from 80 to 8000.

Type **`sudo nano /etc/apache2/ports.conf`**

Change line Listen from 80 to 8000, like in the picture.

Press Ctrl + O to save, Ctrl + X to exit nano.

Reboot Raspberry Pi.

Do this for all Raspberry Pis.

After this, you need to supply the port to access the web server. Like in the picture.

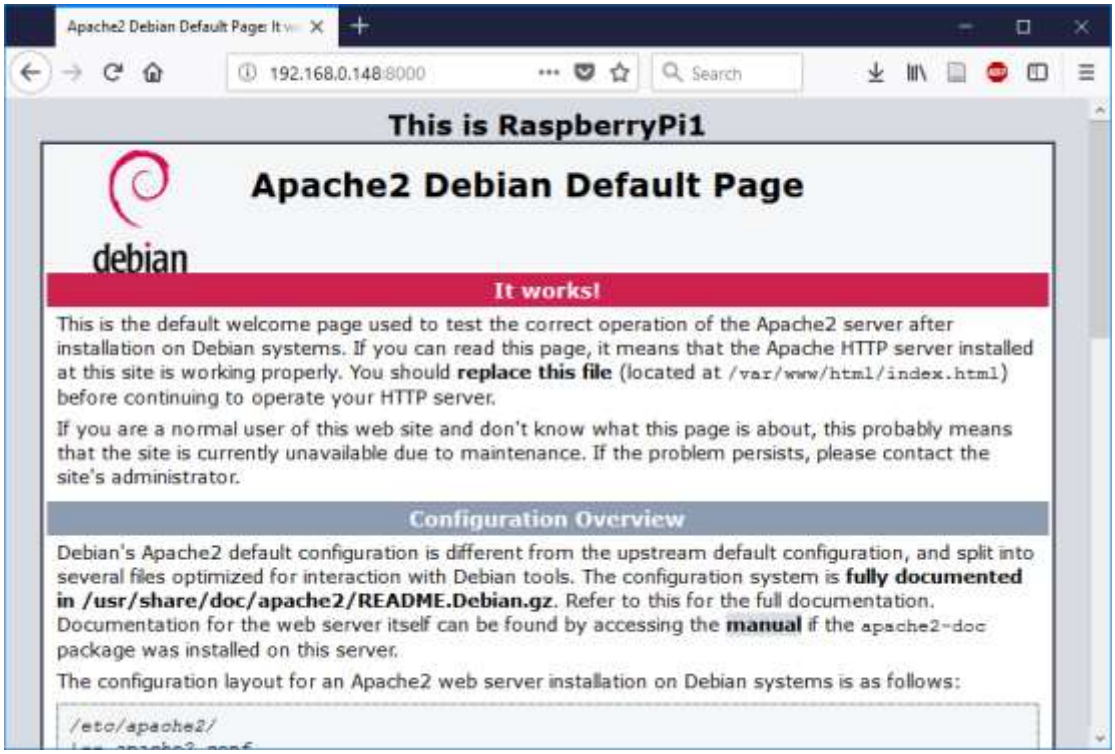
Apache2 Debian Default Page: It
192.168.0.148:8000
Search
This is RaspberryPi1
Apache2 Debian Default Page
debian
It works!
This is the default welcome page used to test the correct operation of the Apache2 server after installation on Debian systems. If you can read this page, it means that the Apache HTTP server installed at this site is working properly. You should replace this file (located at /var/www/html/index.html) before continuing to operate your HTTP server.
If you are a normal user of this web site and don't know what this page is about, this probably means that the site is currently unavailable due to maintenance. If the problem persists, please contact the site's administrator.
Configuration Overview
Debian's Apache2 default configuration is different from the upstream default configuration, and split into several files optimized for interaction with Debian tools. The configuration system is fully documented in /usr/share/doc/apache2/README.Debian.gz. Refer to this for the full documentation. Documentation for the web server itself can be found by accessing the manual if the apache2-doc package was installed on this server.
The configuration layout for an Apache2 web server installation on Debian systems is as follows:
/etc/apache2/

Part 5: Load balancing

Open a terminal on the first Raspberry Pi.

Update system and install HAPROXY load balancer.

sudo apt-get update

sudo apt-get install haproxy

Edit HAPROXY configuration

sudo nano /etc/haproxy/haproxy.cfg

At the end of the file add the following lines, like in the picture.

```
frontend clusternodes
    bind *:80
    mode http
    default_backend nodes

backend nodes
    mode http
    balance roundrobin
    option forwardfor
    http-request set-header X-Forwarded-Port %[dst_port]
    http-request add-header X-Forwarded-Proto https if { ssl_fc }
    option httpchk HEAD / HTTP/1.1\r\nHost:localhost
    server node1 192.168.0.148:8000 check
    server node2 192.168.0.103:8000 check
    server node3 192.168.0.140:8000  check

listen stats
    bind  *:9000
    stats enable
    stats uri /
    stats hide-version
    stats auth user:pass
```

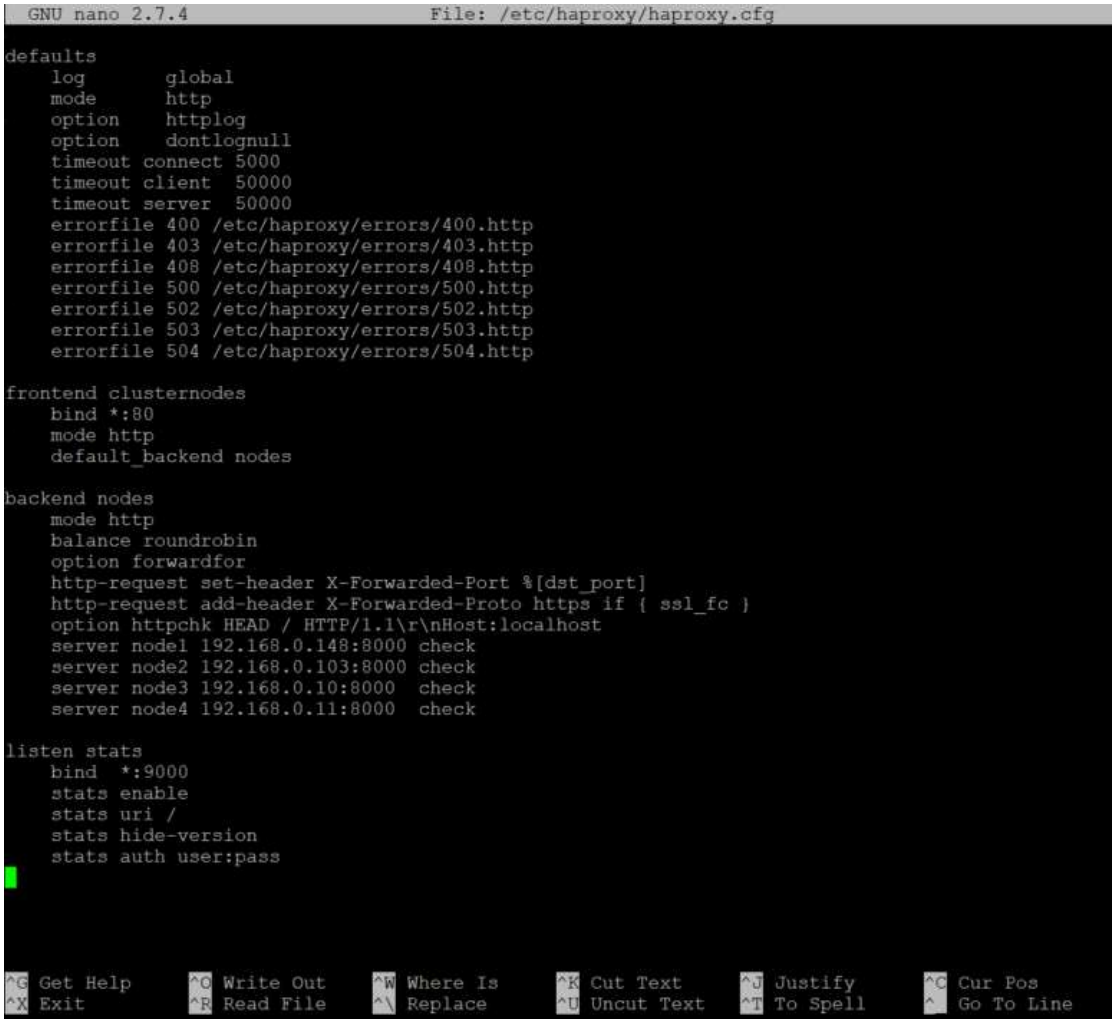

Restart HAPROXY service to apply changes.

```
sudo service haproxy restart
```

Afterwards open a browser and give the IP of node1 (raspberry1). We should see that if you refresh the page continuously, then sometimes it serves the page from Raspberry Pi 1, then from Raspberry Pi 2, and so on. This is load balancing. The workload is distributed among the nodes of the cluster.

If you give the IP of the first Raspberry Pi with the port 9000, then you have access to the statistics of the load balancer. The username is user and the password pass.

These are very useful measurements which must be utilized in the text.

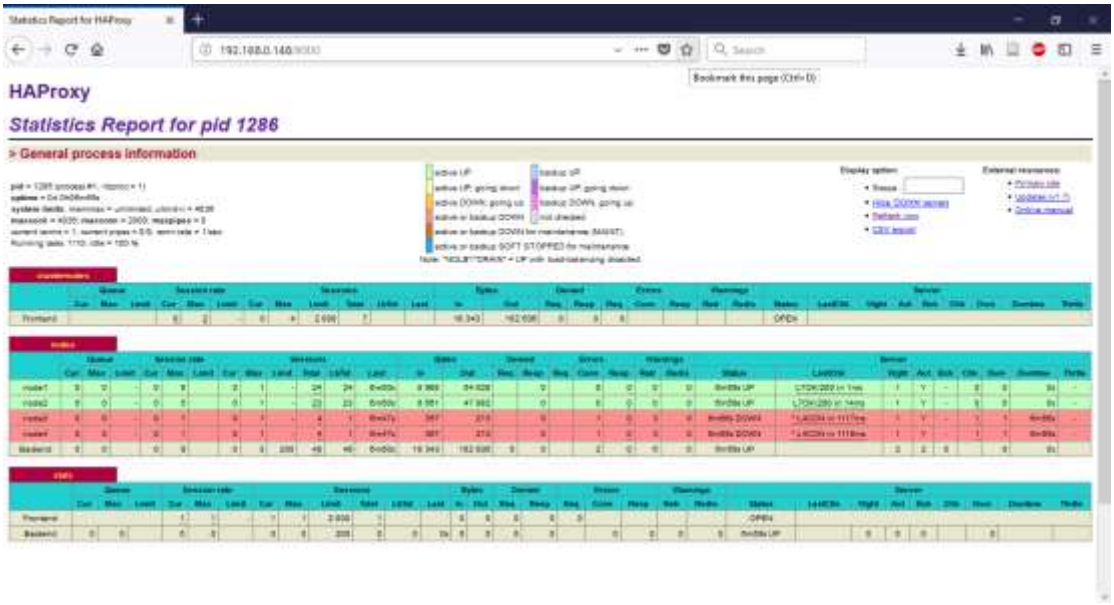

References:

http://www.raspberrywebserver.com/raspberrypicluster/raspberry-pi-cluster.html

http://www.gregtrowbridge.com/setting-up-a-multiple-raspberry-pi-web-server-part-1/

http://www.gregtrowbridge.com/setting-up-a-multiple-raspberry-pi-web-server-part-2/

http://www.gregtrowbridge.com/setting-up-a-multiple-raspberry-pi-web-server-part-3/

http://gregtrowbridge.com/setting-up-a-multiple-raspberry-pi-web-server-part-4/

http://gregtrowbridge.com/setting-up-a-multiple-raspberry-pi-web-server-part-5/

https://serversforhackers.com/c/load-balancing-with-haproxy

https://stackoverflow.com/questions/42348849/configuration-for-haproxy-in-ubuntu